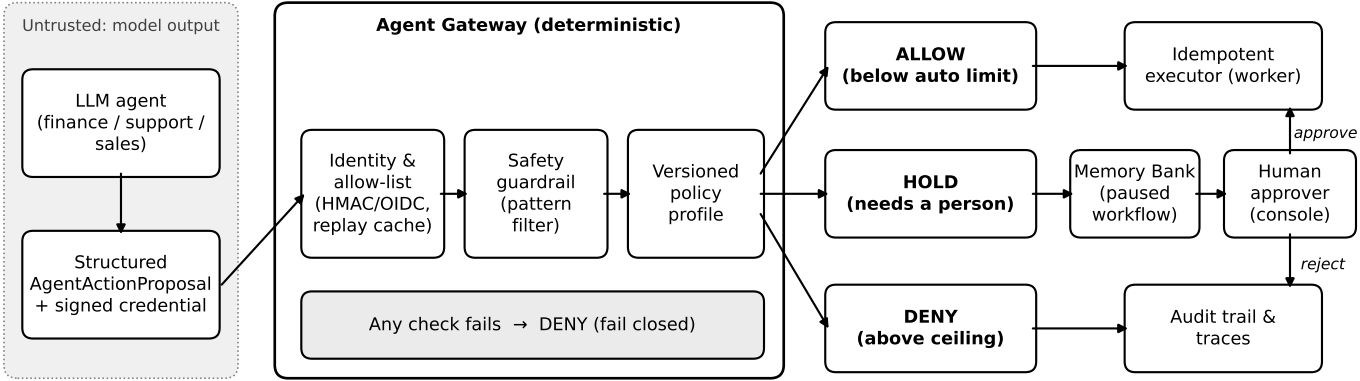
Untrusted: model output
LLM agent (finance / support / sales)
Structured AgentActionProposal + signed credential
Agent Gateway (deterministic)
Identity & allow-list (HMAC/OIDC, replay cache)
Safety guardrail (pattern filter)
Versioned policy profile
Any check fails → DENY (fail closed)
ALLOW (below auto limit)
HOLD (needs a person)
DENY (above ceiling)
Idempotent executor (worker)
Memory Bank (paused workflow)
Human approver (console)
approve
reject
Audit trail & traces